

step-2

Step-2 shows:

- How degrees of freedom are defined with finite elements
- The *DoFHandler* class
- How DoFs are connected by bilinear forms
- Sparsity patterns of matrices
- How to visualize a sparsity pattern

step-2

Sparsity of system matrices:


- For PDEs, finite element matrices are *always sparse*
- Result of
 - *local* definition of shape functions
 - *locality* of the differential operator

Sparsity is not a coincidence. It is a design choice of the finite element method.

Sparsity can not be overestimated as a factor in the success of the FEM!

step-2

Example: Consider this mesh and bilinear form:



$$\begin{aligned} A_{ij} &= (\nabla \varphi_i, \nabla \varphi_j) \\ &= \int_{\Omega} \nabla \varphi_i \cdot \nabla \varphi_j \, dx \end{aligned}$$


Note: In general we have that

- $A_{00} \neq 0, A_{01} \neq 0, A_{02} \neq 0, A_{06} \neq 0$
- $A_{03} = A_{04} = A_{05} = A_{07} = A_{08} = 0$


The bigger the mesh, the more zeros there are per row!

step-2

Renumbering: The order of enumerating degrees of freedom is arbitrary



VS.



Notes:

- Resulting matrices are just permutations of each other
- Both sparse; most algorithms don't care about ordering

Gaussian elimination

End result: finite element matrices are sparse. So how do we solve linear systems $AU=F$?

Recall:

- Gaussian elimination adds one row to the ones below
- The rows below therefore gain nonzero entries
- What started out as a matrix with cN nonzero entries will result in something with $O(N^2)$ entries!
- In other words, the result is a *full* matrix!

Gaussian elimination

Gaussian elimination leads to a full matrix.

Why is this bad:

- We need to store $O(N^2)$ entries instead of cN entries
- Computing $O(N^2)$ entries generally costs $O(N^3)$ ops

Imagine cases where $N=1,000,000$.

Remark: There are “sparse direct solvers” that try to be smarter. But the basic problem remains.

Iterative solvers

Idea: Do not try to actually factor/invert the matrix A .

Instead: Use methods that

- solve $AU=F$ by starting with a guess for U and *iteratively* improving on it
- only ever *multiply* by A

Benefits:

- $O(N)$ memory consumption
- $O(N)$ cost for each iteration
- $O(N)$ overall cost if we could get away with a fixed number of iterations

Iterative solvers

Idea: Do not try to actually factor/invert the matrix A .

Instead: Use methods that

- solve $AU=F$ by starting with a guess for U and *iteratively* improving on it
- only ever *multiply* by A

Examples of methods:

- Fixed point iterations: Richardson, Gauss-Seidel, SOR, ...
- Krylov subspace methods: Conjugate Gradients (CG), GMRES, BiCGStab, ...

Direct vs iterative

Guidelines for direct solvers vs iterative solvers:

Direct solvers:

- ✓ *Always* work, for any invertible matrix
- ✓ No need to think about preconditioners
- ✓ Faster for problems with $<100k$ unknowns
- ✗ Need too much memory + CPU time for larger problems

Iterative solvers:

- ✓ Need $O(N)$ memory
- ✓ Can solve *very* large problems
- ✓ Often parallelize well
- ✗ Choice of solver/preconditioner depends on problem

Advice for iterative solvers

There is a wide variety of iterative solvers:

- CG: Conjugate gradients
- MinRes: Minimal residuals
- GMRES: Generalized minimal residuals
- F-GMRES: Flexible GMRES
- SymmLQ: Symmetric LQ decomposition
- BiCGStab: Biconjugate gradients stabilized
- QMR: Quasi-minimal residual
- TF-QMR: Transpose-free QMR
- ...

Which solver to choose depends on the properties of the matrix, primarily *symmetry* and *definiteness*!

Advice for iterative solvers

Guidelines for use:

- CG: Matrix is symmetric, positive definite
- MinRes: –
- GMRES: Catch-all
- F-GMRES: Catch-all with variable preconditioners
- SymmLQ: –
- BiCGStab: Matrix is non-symmetric but positive definite
- QMR: –
- TF-QMR: –
- All others: –

In reality, only CG, BiCGStab and (F-)GMRES are used much.

Advice for iterative solvers

Note:

**All iterative solvers are bad
without a good preconditioner!**

**The art of devising a good iterative solver
is to devise a good preconditioner!**

Finite element methods in scientific computing

Wolfgang Bangerth, Colorado State University

Lecture 35:

What preconditioner to use

Introduction

Parts 1+2: Simple preconditioners for simple problems

Observations on iterative solvers

The finite element method provides us with a linear system

$$Ax = b$$

that we then need to solve.

Basic observations:

- For sparse direct solvers, speed of solution only depends on sparsity pattern
- For iterative solvers, performance also depends on the *values* in A
- Performance measures:
 - number of iterations
 - cost of every iteration

Observations on iterative solvers

The finite element method provides us with a linear system

$$Ax = b$$

that we then need to solve.

Factors affecting performance of iterative solvers:

- Symmetry of a matrix
- Whether A is definite
- Condition number of A
- How the eigenvalues of A are clustered
- Whether A is reducible/irreducible

Observations on iterative solvers

Example 1: Using CG to solve

$$Ax = b$$

where A is SPD, each iteration reduces the residual by a factor of

$$r = \frac{\sqrt{\kappa(A)} - 1}{\sqrt{\kappa(A)} + 1} < 1$$

- For a tolerance ϵ we need $n = \frac{\log \epsilon}{\log r}$ iterations

- For the Laplace matrix, this is $r = \frac{c-h}{c+h} < 1$
 $n = O\left(\frac{\log \epsilon}{h}\right)$

Observations on iterative solvers

Example 2: When solving

$$Ax = b$$

where A has the form

$$A = \begin{pmatrix} a_{11} & 0 & 0 & \cdots \\ 0 & a_{22} & 0 & \cdots \\ 0 & 0 & a_{33} & \cdots \\ \vdots & \vdots & \vdots & \ddots \end{pmatrix}$$

then every decent iterative solver converges in 1 iteration.

Note 1: This, even though condition number may be large

Note 2: This is true, in particular, if $A=I$.

The idea of preconditioners

Idea: When solving

$$Ax = b$$

maybe we can find a matrix P^{-1} and instead solve

$$P^{-1}Ax = P^{-1}b$$

Observation 1: If $P^{-1}A \approx D$ then solving should require fewer iterations

Corollary: The perfect preconditioner is a multiple of the inverse matrix, i.e., $P^{-1} = A^{-1}$.

The idea of preconditioners

Idea: When solving

$$Ax = b$$

maybe we can find a matrix P^{-1} and instead solve

$$P^{-1}Ax = P^{-1}b$$

Observation 2: Iterative solvers only need matrix-vector multiplications, no element-by-element access.

The idea of preconditioners

Idea: When solving

$$Ax = b$$

maybe we can find a matrix P^{-1} and instead solve

$$P^{-1}Ax = P^{-1}b$$

Observation 3: There is a tradeoff:
fewer iterations vs cost of preconditioner.

Corollary: Preconditioning only works if P^{-1} is cheap to compute and if P^{-1} is cheap to *apply* to a vector.

Example: $P^{-1}=A^{-1}$ does not qualify.

Constructing preconditioners

Part 1: Constructing simple preconditioners for elliptic problems (1952 - 1980s)

Constructing preconditioners

Remember: When solving the preconditioned system

$$P^{-1}Ax = P^{-1}b$$

then the best preconditioner is $P^{-1}=A^{-1}$.

Problem: (i) We can't compute it efficiently. (ii) If we could, we would not need an iterative solver.

But: Maybe we can approximate $P^{-1} \approx A^{-1}$.

Idea 1: Do we know of other iterative solution techniques?

Idea 2: Use incomplete decompositions.

Constructing preconditioners

Approach 1: Remember the oldest iterative techniques!

To solve $Ax = b$ we can use *defect correction*:

- Under certain conditions, the iteration:

$$x^{(k+1)} = x^{(k)} - P^{-1}(Ax^{(k)} - b)$$

will converge to the exact solution x

- Unlike Krylov-space methods, convergence is linear
- The best preconditioner is again $P^{-1} \approx A^{-1}$

Constructing preconditioners

Approach 1: Remember the oldest iterative techniques!

Preconditioned defect correction for $Ax = b$, $A = L+D+U$:

- Jacobi iteration:

$$x^{(k+1)} = x^{(k)} - \omega D^{-1}(Ax^{(k)} - b)$$

- The Jacobi preconditioner is then

$$P^{-1} = \omega D^{-1}$$

which is easy to compute and apply.

Note: We don't need the scaling (“relaxation”) factor.

Constructing preconditioners

Approach 1: Remember the oldest iterative techniques!

Preconditioned defect correction for $Ax = b$, $A = L+D+U$:

- Gauss-Seidel iteration:

$$x^{(k+1)} = x^{(k)} - \omega (L+D)^{-1} (Ax^{(k)} - b)$$

- The Gauss-Seidel preconditioner is then

$$P^{-1} = \omega (L+D)^{-1} \quad \text{i.e. } h = P^{-1}r \text{ solves } (L+D)h = \omega r$$

which is easy to compute and apply as $L+D$ is triangular.

Note 1: We don't need the scaling (“relaxation”) factor.

Note 2: This preconditioner is not symmetric.

Constructing preconditioners

Approach 1: Remember the oldest iterative techniques!

Preconditioned defect correction for $Ax = b$, $A = L+D+U$:

- SOR (Successive Over-Relaxation) iteration:

$$x^{(k+1)} = x^{(k)} - \omega (D + \omega L)^{-1} (Ax^{(k)} - b)$$

- The SOR preconditioner is then

$$P^{-1} = (D + \omega L)^{-1}$$

Note: This preconditioner is not symmetric.

Constructing preconditioners

Approach 1: Remember the oldest iterative techniques!

Preconditioned defect correction for $Ax = b$, $A = L+D+U$:

- SSOR (Symmetric Successive Over-Relaxation) iteration:

$$x^{(k+1)} = x^{(k)} - \frac{1}{\omega(2-\omega)} (D+\omega U)^{-1} D (D+\omega L)^{-1} (Ax^{(k)} - b)$$

- The SSOR preconditioner is then

$$P^{-1} = (D+\omega U)^{-1} D (D+\omega L)^{-1}$$

Note: This preconditioner is now symmetric if A is symmetric!

Constructing preconditioners

Approach 1: Remember the oldest iterative techniques!

Common observations about preconditioners from stationary iterations:

- Have been around for a long time
- Generally useful for small problems ($<100,000$ DoFs)
- Not particularly useful for larger problems

Constructing preconditioners

Approach 2: Approximations to A^{-1}

Idea 1: Incomplete decompositions

- Incomplete LU (ILU):
Perform an LU decomposition on A but only keep elements of L , U that fit into the sparsity pattern of A
- Incomplete Cholesky (IC):
 LL^T decomposition if A is symmetric
- Many variants:
 - strengthen diagonal
 - augment sparsity pattern
 - thresholding of small/large elements

Summary

Preconditioners for “simple” problem:

- Defect correction-based preconditioners are the simplest choice
- Limited by slow speed of information propagation
- ILU/IC frequently better but more complex and limited by memory requirements/CPU time
- Both kinds work reasonably well for “small” problems
- For elliptic problems we today have much better methods: *Geometric* and *Algebraic Multigrid (GMG, AMG)*

Overall summary

Sparsity has many implications:

- Finite element matrices
 - can be stored efficiently
 - can be multiplied by efficiently
 - can not be inverted efficiently
- Iterative methods work around that:
 - only require multiplication
 - iteratively improve the solution
- **But:** Speed of convergence depends on properties of A
- “Preconditioners” address this by solving $P^{-1}A U = P^{-1}F$ instead of $A U = F$.